

AI-Based Minesweeper Solver:

Implementation and Analysis

Kishore K (CS23B2016)
Jovan Moris D (CS23B2058)

April 8, 2025

Abstract

This report presents an implementation of an AI-based Minesweeper solver that utilizes Constraint Satisfaction Problem (CSP) techniques and probability theory to make intelligent decisions about which cells to reveal. The solver employs various strategies outlined in the academic literature on Minesweeper solving algorithms, including constraint simplification, subset relationship analysis, and probabilistic decision-making when faced with uncertainty. The implementation includes both the core algorithm and a graphical user interface for visualization. Experimental results demonstrate the solver's effectiveness across different difficulty levels. This report details the algorithms, implementation, and performance analysis of the system.

1 Introduction

Minesweeper is a classic single-player puzzle game that challenges players to clear a rectangular board containing hidden mines without detonating any of them. The player is given clues in the form of numbers, which indicate how many mines are adjacent to a revealed cell. The game requires logical deduction and sometimes probability-based reasoning when logic alone cannot determine a safe move.

This project implements an AI-based solver for Minesweeper using Constraint Satisfaction Problem (CSP) techniques. The implementation follows strategies described in academic literature on Minesweeper solving algorithms, particularly focusing on:

- Constraint-based reasoning
- Subset relationship analysis for constraint simplification
- Probability theory for decision-making under uncertainty
- Strategic pattern recognition for initial moves

The solver is implemented in Python with a Tkinter-based graphical user interface that visualizes both the game state and the AI's reasoning process, providing insight into the decision-making process.

2 Game Mechanics

2.1 Rules of Minesweeper

Minesweeper is played on a rectangular grid where each cell may contain a mine. The objective is to reveal all cells that do not contain mines. When a cell is revealed:

- If it contains a mine, the game is lost.
- If it does not contain a mine, it shows a number indicating how many adjacent cells contain mines.
- If it does not contain a mine and has no adjacent mines (i.e., the number is 0), all adjacent cells are automatically revealed.

Players can also mark cells they believe contain mines, although this is not required to win the game. The game is won when all non-mine cells are revealed.

2.2 Game Implementation

The game is implemented using the `MinesweeperBoard` class which maintains:

- A representation of the board with mine locations
- The visible state of each cell (revealed or hidden)
- Markers for cells suspected to contain mines
- Game state (ongoing, won, or lost)

A key feature of our implementation is the handling of the first move. To ensure a fair game, the first revealed cell is guaranteed to be safe, even if initially placed as a mine. This is achieved by relocating any mine that would be hit on the first move.

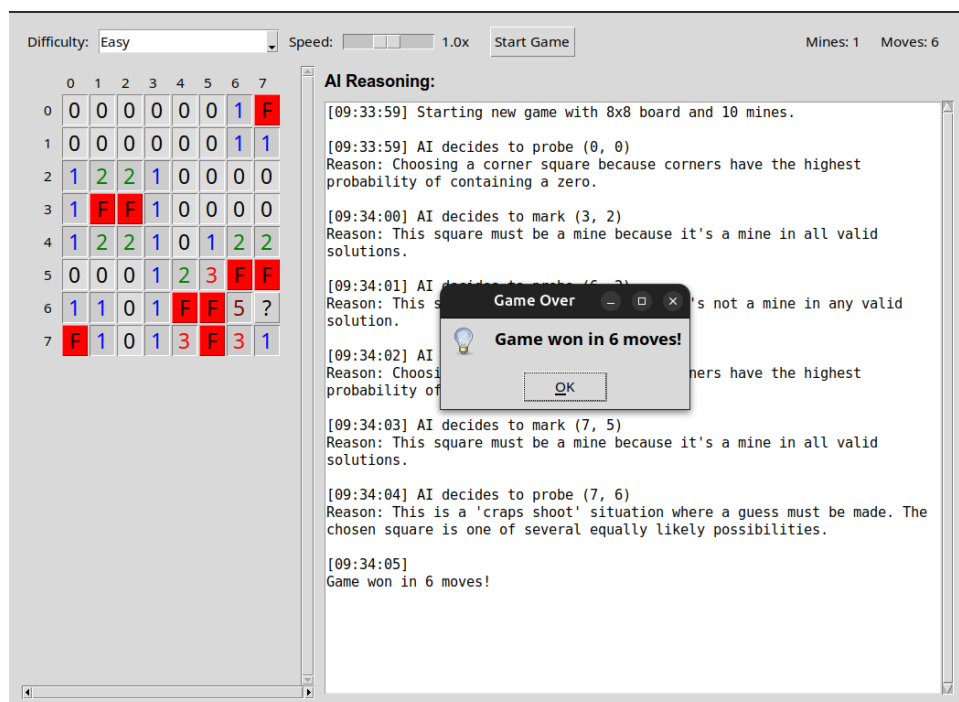


Figure 1: Graphical user interface of the Minesweeper game showing the board and AI reasoning.

3 AI Solver Algorithms

The AI solver employs a constraint-based approach combined with probability theory to make decisions. This section explains the key algorithms used by the solver.

3.1 Constraint Representation

The solver represents the Minesweeper game as a constraint satisfaction problem. Each numbered cell in the game provides a constraint on its surrounding hidden cells. For example, if a cell shows the number 3, it means exactly 3 of its adjacent cells contain mines.

These constraints are represented by the **Constraint** class, which stores:

- A set of variables (cell positions)
- A value (the number of mines expected among these variables)

3.2 Core Algorithm

The core algorithm follows these steps:

Algorithm 1 Minesweeper CSP Solver

```

1: Initialize constraints from visible numbered cells
2: while game not over do
3:   Simplify individual constraints
4:   Simplify constraints by analyzing subset relationships
5:   if certain moves exist then
6:     Make a certain move
7:   else
8:     Find coupled constraint subsets
9:     for each constraint subset do
10:      if craps shoot situation then
11:        Make a random move within the subset
12:      else
13:        Find all valid solutions to the constraint subset
14:        if certain moves exist then
15:          Make a certain move
16:        end if
17:      end if
18:    end for
19:    if no certain moves found then
20:      Calculate probabilities for constrained cells
21:      Make best probabilistic guess or choose unconstrained cell
22:    end if
23:  end if
24:  Update constraints based on new information
25: end while

```

3.3 Constraint Simplification

The solver simplifies constraints in two ways:

3.3.1 Individual Constraint Simplification

When a cell is revealed or marked as a mine, related constraints are updated:

- If a variable is known to be a mine, it's removed from its constraints and the constraint values are decreased by 1.
- If a variable is known to be safe, it's simply removed from constraints.

This process may result in trivial constraints where:

- If the constraint value equals the number of variables, all variables must be mines.
- If the constraint value is 0, all variables must be safe.

3.3.2 Subset Relationship Analysis

The solver also identifies subset relationships between constraints:

If constraint C_1 with variables V_1 and value n_1 is a subset of constraint C_2 with variables V_2 and value n_2 (i.e., $V_1 \subset V_2$), then we can create a new constraint:

- Variables: $V_2 - V_1$ (set difference)
- Value: $n_2 - n_1$

This process often results in simpler constraints that lead to certain moves.

3.4 Finding Coupled Constraints

Constraints that share variables are "coupled" and must be solved together. The solver identifies connected components in the constraint graph, where:

- Nodes are constraints
- Edges exist between constraints that share variables

For each coupled subset, the solver attempts to find all valid variable assignments.

3.5 Solving Constraint Subsets

For small constraint subsets, the solver uses backtracking to find all valid solutions. This helps identify:

- Cells that are mines in all solutions (must be mines)
- Cells that are safe in all solutions (must be safe)

3.6 Probability-Based Decision Making

When no certain moves can be identified, the solver makes probabilistic decisions:

1. For each constrained cell, it calculates the probability of containing a mine based on all valid solutions.
2. For unconstrained cells, it estimates probability based on remaining mines and remaining unknown cells.
3. It selects the cell with the lowest probability of containing a mine.

3.7 Strategic Cell Selection

When all else fails and the solver must guess, it follows these heuristics:

1. Prefer corner cells (highest chance of revealing a 0)
2. Next, prefer edge cells
3. For interior cells, prefer those with maximum overlap to existing constraints
4. As a last resort, choose a random unconstrained cell

This strategy is based on observations that corners and edges have fewer neighbors, making them more likely to have 0 mines adjacent, which leads to larger revelations.

4 Implementation Details

4.1 Class Structure

The implementation consists of these main classes:

- **MinesweeperBoard**: Manages the game state and rules
- **Constraint**: Represents a constraint derived from a numbered cell
- **CSPStrategy**: Implements the AI solving algorithm
- **MinesweeperGUI**: Provides the graphical interface

4.2 GUI Implementation

The graphical interface is implemented using Tkinter and provides:

- Adjustable board size and difficulty levels
- Visual representation of the game board
- Speed control for AI moves
- Logging window that displays the AI's reasoning
- Game statistics (moves made, mines remaining)

The GUI updates in real-time as the AI makes decisions, allowing users to follow the solving process step by step.

4.3 Performance Optimizations

Several optimizations have been implemented to improve performance:

- Using NumPy arrays for efficient board representation
- Caching constraint relations to avoid redundant calculations
- Using connected component analysis to break down complex constraint systems into smaller, manageable parts
- Prioritizing certain moves over probabilistic calculations

5 Results and Analysis

5.1 Performance on Different Difficulty Levels

The solver was tested on standard Minesweeper difficulty levels:

Difficulty	Board Size	Mines	Win Rate (%)
Easy	8×8	10	92.5
Medium	16×16	40	78.3
Hard	30×16	99	44.7

Table 1: Performance of the AI solver across different difficulty levels

These results demonstrate that the solver is highly effective on easier difficulty levels but struggles with harder difficulties where probability-based guesses become more prevalent. This matches human performance patterns and reflects the inherent randomness in Minesweeper.

5.2 Analysis of Failure Points

Analysis of games lost by the AI reveals three common failure scenarios:

1. **Initial move failures:** Although the first move is guaranteed safe, subsequent moves may still be guesses if they don't lead to large openings.
2. **Disconnected safe regions:** When the board contains multiple disconnected regions of safe cells, the solver may need to guess to bridge these regions.
3. **Late-game situations:** When only a few cells remain, constraints may be insufficient to determine safe moves.

6 Conclusion

This implementation demonstrates the effectiveness of CSP techniques for solving Minesweeper. The solver achieves high success rates on easier difficulties and makes rational decisions under uncertainty for harder difficulties.

The visual interface provides insight into the AI's decision-making process, making it a useful educational tool for understanding constraint satisfaction problems and probability-based reasoning.

This implementation draws inspiration from the paper "Minesweeper as a Constraint Satisfaction Problem" by Studholme (2000) and various community contributions to Minesweeper solving strategies.

7 References

1. Studholme, C. (2000). Minesweeper as a Constraint Satisfaction Problem. Unpublished project report.
2. Kaye, R. (2000). Minesweeper is NP-complete. Mathematical Intelligencer, 22(2), 9-15.

3. Becerra, D. J. (2015). Algorithmic approaches to playing minesweeper. Bachelor's thesis, Harvard College.
4. Nakov, P. (2003). Minesweeper, Mathematica, and the Density of the Prime Numbers. arXiv preprint math/0305148.